
Using GLRM on H2O.ai

— *to Predict Customer Value* —
Corinne Powers, PhD

About Me

Sr. Data Scientist at Retina AI

Math, Optimization and Signals Modeling Background

R and Python

New to LA Community (8 months)



RETINA

This Talk

Business Value

Modeling & Data

R Usage in H2O.ai

Motivation



- You sell Animal Cookies through your website
- Customers may buy animal packs or subscribe monthly
- You track session data of site visits, even if they don't purchase

How do we **infer** behavior to **target** customers who could be valuable?

Clustering!

Big Picture: Seek people who “look” similar to high-value customers.

Supervised Learning

Decision tree

K-nearest neighbors

Unsupervised Learning

K-Means clustering

But where can this fail? Any why?

Mixed Data Types

Numerical

Categorical

Boolean

	<i>Total Spend</i>	<i>Device Type</i>	<i>Acquisition Channel</i>	<i>Favorite Shape</i>	<i>Favorite Flavor</i>	<i>Subscriber</i>
Customer A	\$\$	Mobile	Facebook	Tiger	Sprinkles	No
Customer B	\$	Mobile	Google	Elephant	Chocolate	No
Customer C	\$\$\$\$\$	Desktop	Bing	Lion	Sprinkles	Yes
Customer D	\$\$\$	Tablet	Google	Elephant	Sprinkles	No

Semi Supervised Learning

Numerical

Categorical

Boolean

	<i>Total Spend</i>	<i>Device Type</i>	<i>Acquisition Channel</i>	<i>Favorite Shape</i>	<i>Favorite Flavor</i>	<i>Subscriber</i>
Customer A	\$\$	Mobile	Facebook	Tiger	Sprinkles	No
Customer B	\$	Mobile	Google	Elephant	Chocolate	No
Customer C	\$\$\$\$\$	Desktop	Bing	Lion	Sprinkles	Yes
Customer D		Tablet	Google	Elephant	Sprinkles	

Missing Data

Numerical

Categorical

Boolean

	<i>Total Spend</i>	<i>Device Type</i>	<i>Acquisition Channel</i>	<i>Favorite Shape</i>	<i>Favorite Flavor</i>	<i>Subscriber</i>
Customer A	\$\$	Mobile	Facebook	Tiger	Sprinkles	No
Customer B	\$	Mobile	Google			No
Customer C	\$\$\$\$\$	Desktop	Bing	Lion		Yes
Customer D		Tablet	Google	Elephant	Sprinkles	

Generalized Low Rank Models (GLRMs)


Capabilities: handle missing data, handle weird types of data, you pick feature importance

Special Cases: SVM, K-means clustering, PCA

Business Value: Build Personas and Predict \$\$\$, Consumer Preferences, etc.

<https://web.stanford.edu/~boyd/papers/pdf/glrm.pdf>

How it works



<http://docs.h2o.ai/h2o/latest-stable/h2o-docs/data-science/qlrm.html>

GLRM parameters on H2O.ai

```
1 library(h2o)
2 h2o.init()
```

```
1 # GLRM parameters
2 glrm_loss <- "Quadratic"      # Least squares
3 glrm_regx <- "UnitOneSparse"  # Assign to one cluster
4 glrm_gammax <- 1.0
5 glrm_gammay <- 1.0
6 glrm_regy <- "NonNegative"    # Interpretable cluster
7 glrm_k <- 8                   # Number of clusters
```

<https://cran.r-project.org/web/packages/h2o/index.html>

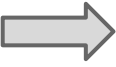
Train a GLRM

```
1 visits.glm <- h2o.glm(  
2   training_frame = visits.train,  
3   cols = union(glm_features, glm_responses),  
4   k = glm_k,  
5   loss = glm_loss,  
6   regularization_x = glm_regx, regularization_y = glm_regy,  
7   gamma_x = glm_gammax, gamma_y = glm_gammay,  
8   transform = "NONE",  
9   max_iterations = 1000, max_updates = 2000,  
10  svd_method = "GramSVD", min_step_size = 1e-6,  
11  ignore_const_cols = FALSE, seed = 10)
```

Impute Labels via Clusters

```
1 # impute missing values
2 visits.predict <- h2o.reconstruct(visits.glm,
3                                   visits.train[,union(glm_features, glm_responses)])
```

	Arch1	Arch2	Arch3	Arch4	Arch5	Arch6	Arch7	Arch8		reconstr_revenue
[1,]	0	0	0	0	0	0	1	0	1	104.7450
[2,]	1	0	0	0	0	0	0	0	2	110.5066
[3,]	0	0	0	0	1	0	0	0	3	108.6246
[4,]	0	0	0	0	1	0	0	0	4	108.6246
[5,]	0	0	0	0	0	0	1	0	5	104.7450



Target Valuable Customers



- GLRM to create Personas that are **descriptive** and **predictive**
- Match customers to closest Persona to make behavioral inferences
- Even with little data, can match using subset of features as needed

Acknowledgements

Adam Brownell, Emad Hasan, Brad Ito (Retina AI)

- Craft Customer Value business cases
- Support an H2O.ai environment in daily workflow
- Feedback from earlier talks

Anqi Fu (Stanford University)

- Work on GLRM tools for H2O.ai in R

Questions?

corinne@retina.ai

GLRM paper: <https://web.stanford.edu/~boyd/papers/pdf/glrm.pdf>

H2O.ai Cran: <https://cran.r-project.org/web/packages/h2o/index.html>

GLRM on H2O.ai: <http://docs.h2o.ai/h2o/latest-stable/h2o-docs/data-science/glrm.html>