

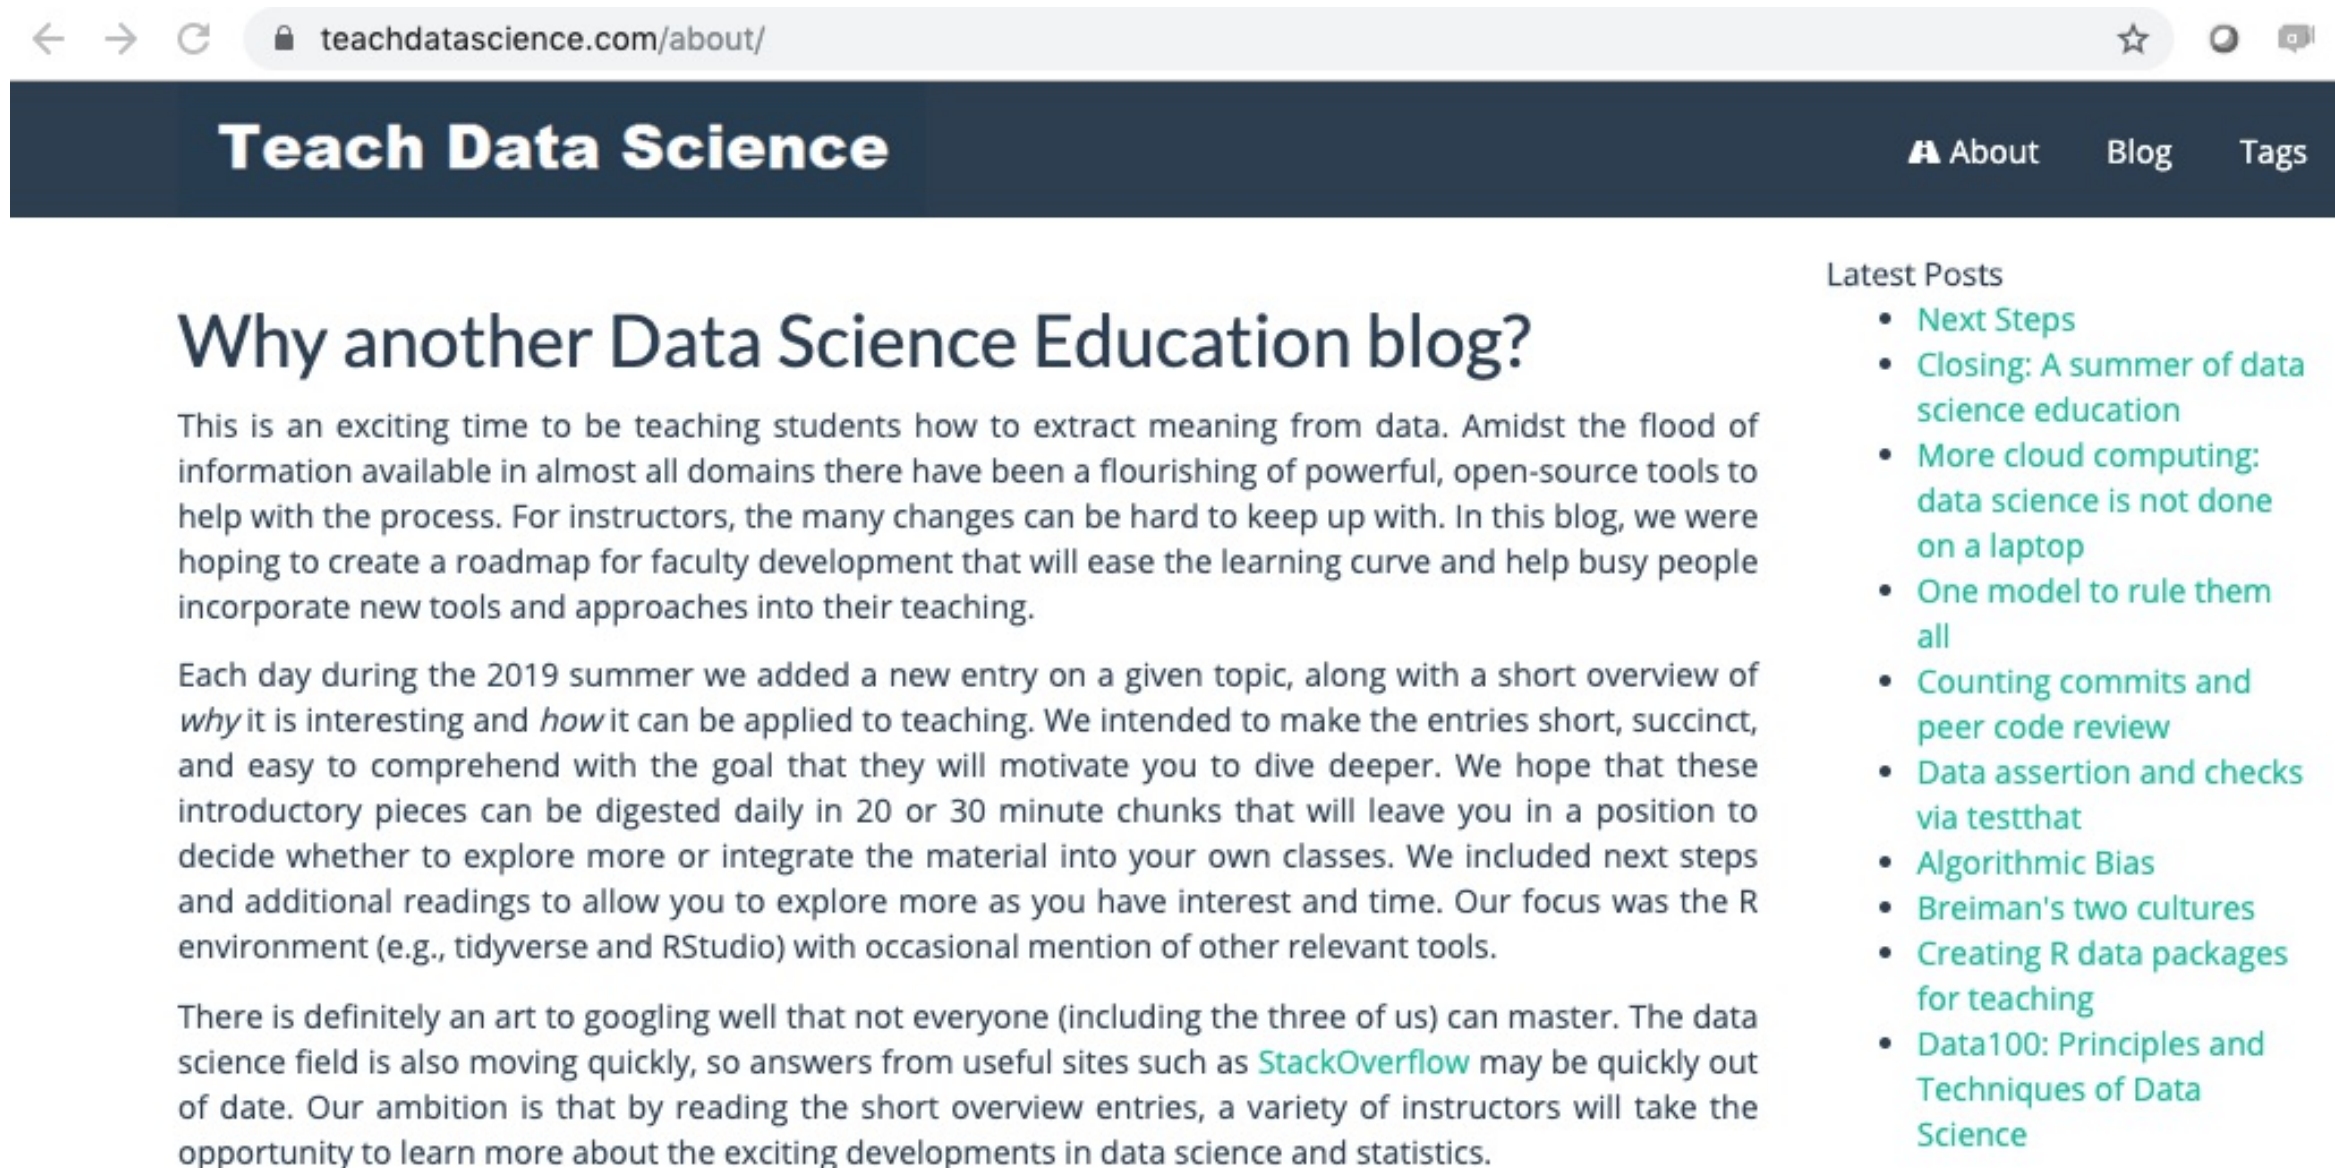
reticulate:

running Python within RStudio

Jo Hardin

October 14, 2019

Just say Yes! -- <https://teachdatascience.com/>



The screenshot shows a web browser window with the address bar displaying "teachdatascience.com/about/". The website has a dark blue header with the title "Teach Data Science" in white. To the right of the title are navigation links: "About", "Blog", and "Tags". The main content area has a large heading "Why another Data Science Education blog?". Below this heading are three paragraphs of text. To the right of the main content is a sidebar titled "Latest Posts" containing a list of ten blog post titles, each preceded by a bullet point.

← → ↻ 🔒 teachdatascience.com/about/ ☆ 🔍

Teach Data Science

About Blog Tags

Why another Data Science Education blog?

This is an exciting time to be teaching students how to extract meaning from data. Amidst the flood of information available in almost all domains there have been a flourishing of powerful, open-source tools to help with the process. For instructors, the many changes can be hard to keep up with. In this blog, we were hoping to create a roadmap for faculty development that will ease the learning curve and help busy people incorporate new tools and approaches into their teaching.

Each day during the 2019 summer we added a new entry on a given topic, along with a short overview of *why* it is interesting and *how* it can be applied to teaching. We intended to make the entries short, succinct, and easy to comprehend with the goal that they will motivate you to dive deeper. We hope that these introductory pieces can be digested daily in 20 or 30 minute chunks that will leave you in a position to decide whether to explore more or integrate the material into your own classes. We included next steps and additional readings to allow you to explore more as you have interest and time. Our focus was the R environment (e.g., tidyverse and RStudio) with occasional mention of other relevant tools.

There is definitely an art to googling well that not everyone (including the three of us) can master. The data science field is also moving quickly, so answers from useful sites such as [StackOverflow](#) may be quickly out of date. Our ambition is that by reading the short overview entries, a variety of instructors will take the opportunity to learn more about the exciting developments in data science and statistics.

Latest Posts

- [Next Steps](#)
- [Closing: A summer of data science education](#)
- [More cloud computing: data science is not done on a laptop](#)
- [One model to rule them all](#)
- [Counting commits and peer code review](#)
- [Data assertion and checks via testthat](#)
- [Algorithmic Bias](#)
- [Breiman's two cultures](#)
- [Creating R data packages for teaching](#)
- [Data100: Principles and Techniques of Data Science](#)

Teach Data Science: reticulate



teachdatascience.com/reticulate/

Teach Data Science

reticulate: running Python within RStudio

📅 16 Jul, 2019 · by Jo Hardin · Read in about 7 min · (1281 words) · Share this on: [f](#) [t](#) [g+](#) [e](#) [in](#) [m](#)



rmarkdown

reticulate

python

data technologies

data wrangling

jupyterhub

For many statisticians, their go-to software language is R. However, there is no doubt that Python is an equally important language in data science. Indeed, the [Jupyter](#) blog entry from earlier this week described the capacities of writing Python code (as well as R and Julia and other environments) using interactive Jupyter notebooks.

```
knitr::opts_chunk$set(collapse = TRUE)
library(reticulate)
use_virtualenv("r-reticulate")
use_python("F:/Anaconda3", required = TRUE)
py_config()
```

Teaching Python and R

What is reticulate??

(much of the talk taken from: <https://rstudio.github.io/reticulate/>)

R Interface to Python

The **reticulate** package provides a comprehensive set of tools for interoperability between Python and R. The package includes facilities for:

- Calling Python from R in a variety of ways including R Markdown, sourcing Python scripts, importing Python modules, and using Python interactively within an R session.
- Translation between R and Python objects (for example, between R and Pandas data frames, or between R matrices and NumPy arrays).
- Flexible binding to different versions of Python including virtual environments and Conda environments.



Reticulate embeds a Python session within your R session, enabling seamless, high-performance interoperability. If you are an R developer that uses Python for some of your work or a member of data science team that uses both languages, reticulate can dramatically streamline your workflow!

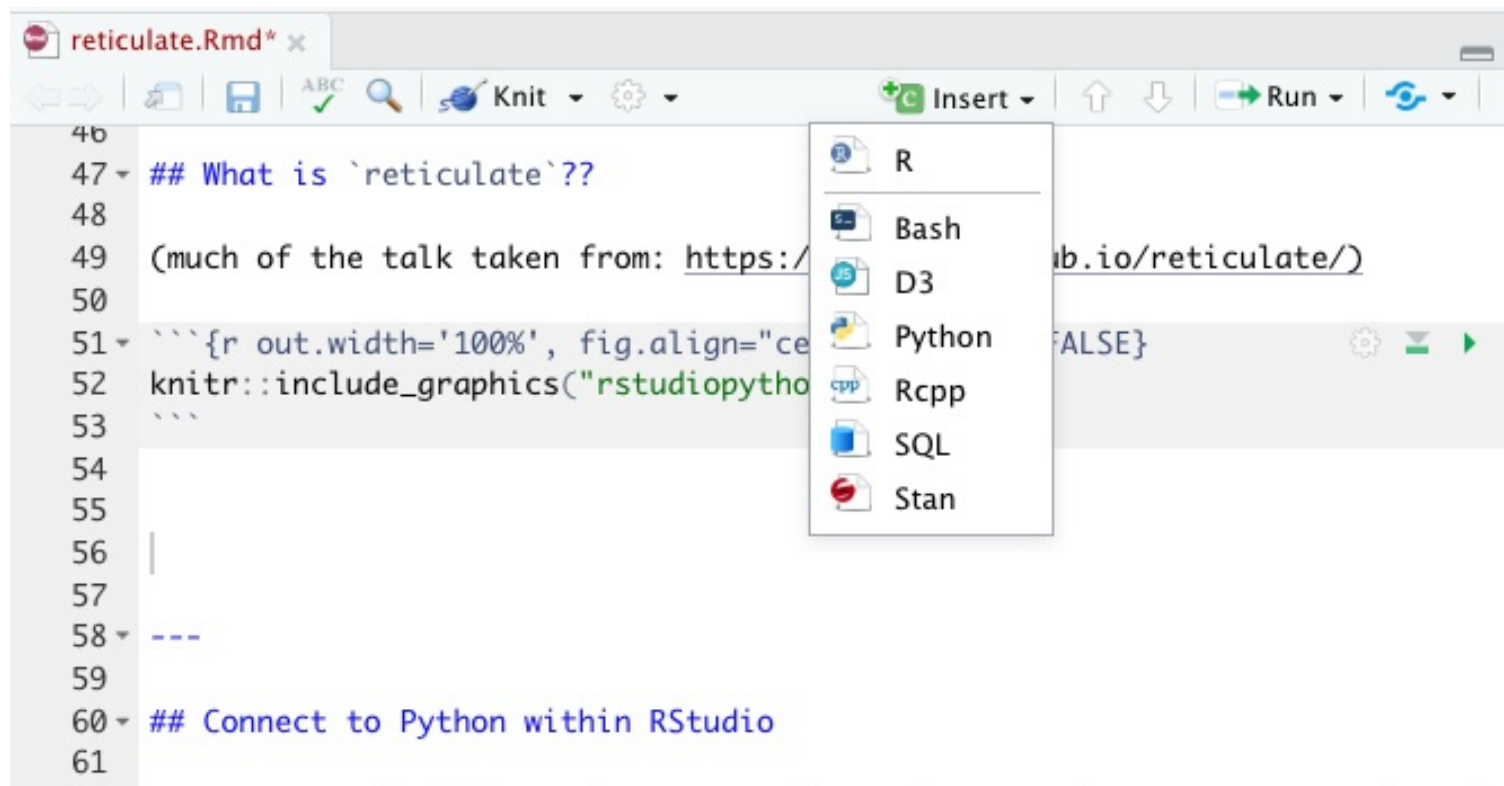
Connect to Python within RStudio

For many statisticians, the go-to software language is R. However, there is no doubt that Python is a very important language in data science. Why not do both??

```
library(reticulate)
use_virtualenv("r-reticulate")
import("statsmodels")
```

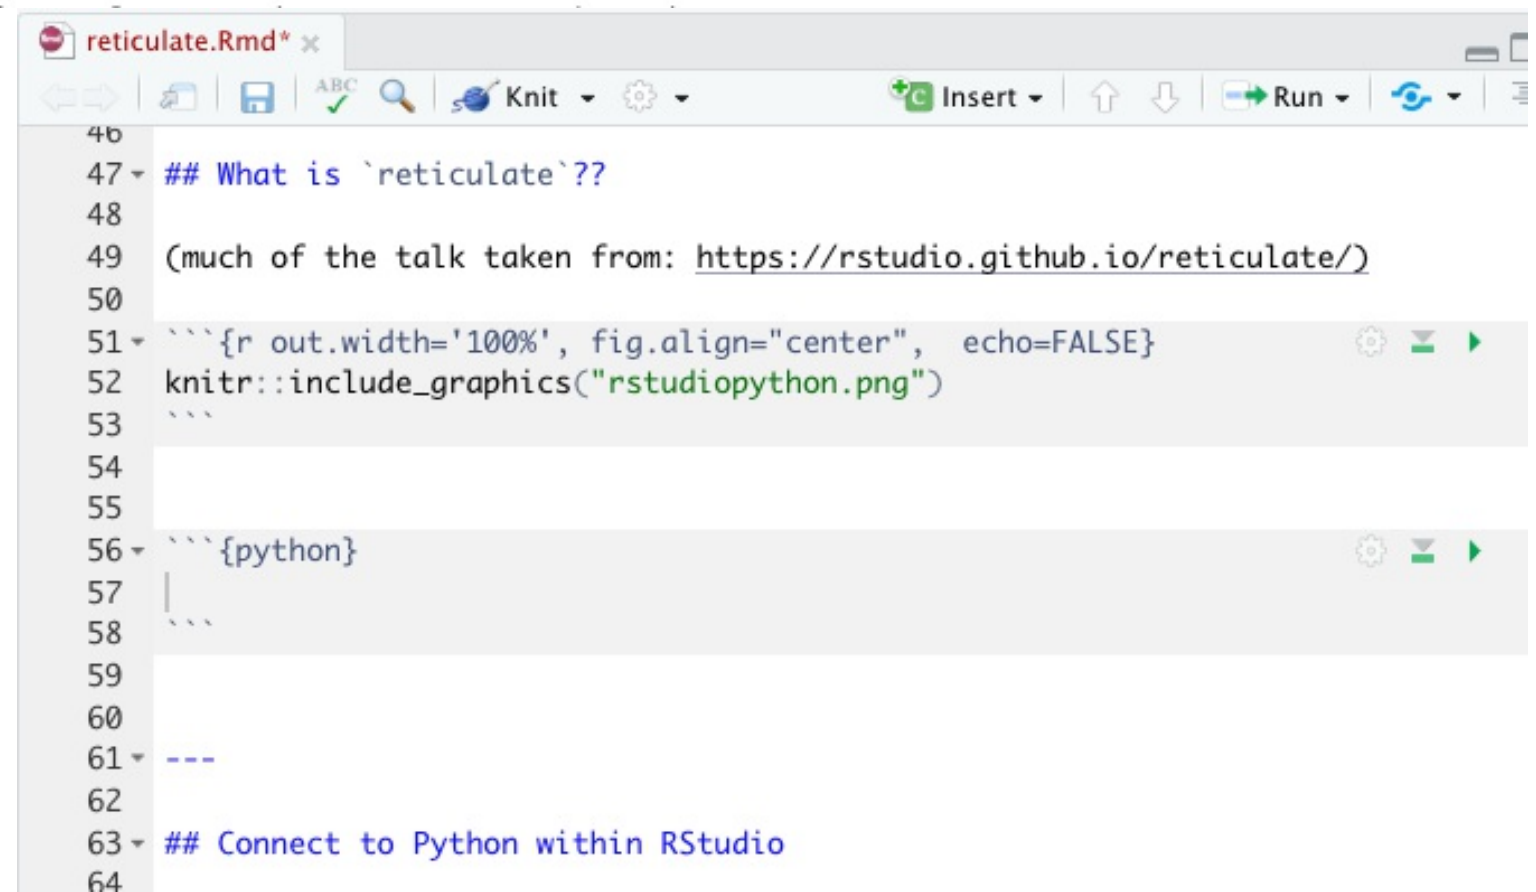
```
## Module(statsmodels)
```

I can run Python inside R??



A screenshot of the RStudio interface. The top toolbar shows the 'Insert' button, which is highlighted. A dropdown menu is open, showing options: R, Bash, D3, Python, Rcpp, SQL, and Stan. The 'Python' option is selected. The background shows a script editor with R code. Line 47 contains the comment '## What is `reticulate`??'. Line 49 contains a URL: '(much of the talk taken from: <https://rstudio.github.io/reticulate/>)'. Line 51 contains a code chunk header: '```{r out.width="100%", fig.align="center", echo=FALSE}'. Line 52 contains the command: 'knitr::include_graphics("rstudiopython.png")'. Line 53 contains the closing backticks: '```'. Line 60 contains the comment '## Connect to Python within RStudio'.

```
46
47 ## What is `reticulate`??
48
49 (much of the talk taken from: https://rstudio.github.io/reticulate/)
50
51 ```{r out.width='100%', fig.align="center", echo=FALSE}
52 knitr::include_graphics("rstudiopython.png")
53 ```
54
55
56
57
58 ---
59
60 ## Connect to Python within RStudio
61
```



A screenshot of the RStudio interface showing the completed R code. The code is the same as in the previous screenshot, but the 'Insert' menu is no longer open. The code is as follows:

```
46
47 ## What is `reticulate`??
48
49 (much of the talk taken from: https://rstudio.github.io/reticulate/)
50
51 ```{r out.width='100%', fig.align="center", echo=FALSE}
52 knitr::include_graphics("rstudiopython.png")
53 ```
54
55
56 ```{python}
57 |
58 ---
59
60 ## Connect to Python within RStudio
61
62
63 ## Connect to Python within RStudio
64
```

Python in R

- pandas for data wrangling.
- In R, the chunk is specified to be a Python chunk (RStudio is now running Python).

```
```{python}
import pandas
flights = pandas.read_csv("flights.csv")
flights = flights[flights["dest"] == "ORD"]
flights = flights[['carrier', 'dep_delay', 'arr_delay']]
flights = flights.dropna()
```

A view of the Python chunk which is actually run:

```
import pandas
flights = pandas.read_csv("flights.csv")
flights = flights[flights["dest"] == "ORD"]
flights = flights[['carrier', 'dep_delay', 'arr_delay']]
flights = flights.dropna()
```

# Learn about the dataset

```
{python}
flights.shape
flights.head(3)
flights.describe()
```

```
flights.shape
```

```
(12590, 3)
```

```
flights.head(3)
```

```
carrier dep_delay arr_delay
4 UA -4.0 12.0
5 AA -2.0 8.0
22 AA -1.0 14.0
```

```
flights.describe()
```

```
dep_delay arr_delay
count 12590.000000 12590.000000
mean 11.709770 2.917951
std 39.409704 44.885155
min -20.000000 -62.000000
25% -6.000000 -22.000000
50% -2.000000 -10.000000
75% 9.000000 10.000000
max 466.000000 448.000000
```



# Computations

```
```{python}
flights = pandas.read_csv("flights.csv")
flights = flights[['carrier', 'dep_delay', 'arr_delay']]
flights.groupby("carrier").mean()
```

```
flights = pandas.read_csv("flights.csv")
flights = flights[['carrier', 'dep_delay', 'arr_delay']]
flights.groupby("carrier").mean()
```

##		dep_delay	arr_delay
##	carrier		
##	AA	8.586016	0.364291
##	AS	5.804775	-9.930889
##	DL	9.264505	1.644341
##	UA	12.106073	3.558011
##	US	3.782418	2.129595

From Python chunk to R chunk


- `py$x` accesses an x variable created within Python from R
- `r.x` accesses an x variable created within R from Python

```
```{r}  
library(ggplot2)
ggplot(py$flights, aes(x=carrier, y = arr_delay)) +
 geom_point() + geom_jitter()
```

# From Python chunk to R chunk

- `py$x` accesses an `x` variable created within Python from R
- `r.x` accesses an `x` variable created within R from Python

```
library(ggplot2)
ggplot(py$flights,
 aes(x=carrier,
 y=arr_delay)) +
 geom_point() +
 geom_jitter()
```



# From R chunk to Python chunk

```
data(diamonds)
head(diamonds)
```

```
A tibble: 6 x 10
carat cut color clarity depth table price x y z
<dbl> <ord> <ord> <ord> <dbl> <dbl> <int> <dbl> <dbl> <dbl>
1 0.23 Ideal E SI2 61.5 55 326 3.95 3.98 2.43
2 0.21 Premium E SI1 59.8 61 326 3.89 3.84 2.31
3 0.23 Good E VS1 56.9 65 327 4.05 4.07 2.31
4 0.290 Premium I VS2 62.4 58 334 4.2 4.23 2.63
5 0.31 Good J SI2 63.3 58 335 4.34 4.35 2.75
6 0.24 Very Good J VVS2 62.8 57 336 3.94 3.96 2.48
```

# A Python chunk

```
print(r.diamonds.describe())
```

```
carat depth ... y z
count 53940.000000 53940.000000 ... 53940.000000 53940.000000
mean 0.797940 61.749405 ... 5.734526 3.538734
std 0.474011 1.432621 ... 1.142135 0.705699
min 0.200000 43.000000 ... 0.000000 0.000000
25% 0.400000 61.000000 ... 4.720000 2.910000
50% 0.700000 61.800000 ... 5.710000 3.530000
75% 1.040000 62.500000 ... 6.540000 4.040000
max 5.010000 79.000000 ... 58.900000 31.800000
##
[8 rows x 7 columns]
```

# A Python chunk

```
import statsmodels.formula.api as smf
model = smf.ols('price ~ carat', data = r.diamonds).fit()
```

```
C:\PROGRA~3\ANACON~1\lib\site-packages\statsmodels\compat\pandas.py:49: FutureWarning: The Panel class is removed from pa
data_klasses = (pandas.Series, pandas.DataFrame, pandas.Panel)
```

```
print(model.summary())
```

```
OLS Regression Results
=====
Dep. Variable: price R-squared: 0.849
Model: OLS Adj. R-squared: 0.849
Method: Least Squares F-statistic: 3.041e+05
Date: Mon, 14 Oct 2019 Prob (F-statistic): 0.00
Time: 10:14:32 Log-Likelihood: -4.7273e+05
No. Observations: 53940 AIC: 9.455e+05
Df Residuals: 53938 BIC: 9.455e+05
Df Model: 1
Covariance Type: nonrobust
=====
coef std err t P>|t| [0.025 0.975]

Intercept -2256.3606 13.055 -172.830 0.000 -2281.949 -2230.772
carat 7756.4256 14.067 551.408 0.000 7728.855 7783.996
=====
Omnibus: 14025.341 Durbin-Watson: 0.986
Prob(Omnibus): 0.000 Jarque-Bera (JB): 153030.525
Skew: 0.939 Prob(JB): 0.00
Kurtosis: 11.035 Cond. No. 3.65
=====
##
Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```



# Bells and whistles

- autocomplete
- Python script

# Importing Python modules

`import()` will import any Python module and call it from R. [The `os` module provides functionality for navigating the operating system.]

```
os <- import("os")
os$listdir(".")
```

```
[1] "custom.css" "flights.csv"
[3] "pychunk1.png" "pychunk2.png"
[5] "PyScript.png" "reticulate.html"
[7] "reticulate.Rmd" "reticulatePyScript.py"
[9] "reticulate_files" "rstudiopython.png"
[11] "ST47Scopy.png" "teachDS.png"
[13] "teachreticulate.png"
```

# Full disclosure

- Python versions (@#\$%#\$%@\$ ????)
- module versions (@%#\$%@\$%#\$ ????)

# Learn more

- RStudio R Interface to Python

<https://rstudio.github.io/reticulate/>

- RStudio blog on Reticulated Python

<https://blog.rstudio.com/2018/10/09/rstudio-1-2-preview-reticulated-python>

# Thank you!

- `jo.hardin@pomona.edu`
- `@jo_hardin47`
- <https://github.com/hardin47>
- <http://research.pomona.edu/johardin/>

